

Generative AI Learning Roadmap for Data Science Engineers

Part 1: Foundations (23 weeks)

- Topics:
 - Generative vs Discriminative models
 - Probability basics: MLE, KL Divergence, Entropy
 - Neural networks refresher (MLP, CNN, RNN)
 - PyTorch or TensorFlow fundamentals
- Resources:
 - Deep Learning Specialization Andrew Ng
 - PyTorch Tutorials

Part 2: Core Generative Models (34 weeks)

- Topics:
 - Variational Autoencoders (VAE)
 - Generative Adversarial Networks (GAN, DCGAN, WGAN, CycleGAN)
 - Autoregressive models (PixelCNN, WaveNet)
- Projects:
 - VAE on MNIST
 - DCGAN on CelebA
- Resources:
 - CS231n (Stanford)
 - GANs in Action (Book)

Part 3: Text & Language Models (34 weeks)

- Topics:
 - Word embeddings (Word2Vec, GloVe)
 - Transformer architecture (Attention, Self-attention)
 - GPT family: GPT-2, GPT-3
- Projects:
 - GPT-2 text generator using HuggingFace

Generative AI Learning Roadmap for Data Science Engineers

- Fine-tune on custom dataset

Part 4: Multimodal Generation (45 weeks)

- Topics:
 - Text-to-image models (Stable Diffusion, DALL·E)
 - Vision Transformers, CLIP
 - Diffusion models
- Projects:
 - Image generation from text using Stable Diffusion
 - CLIP-guided image creation
- Resources:
 - HuggingFace Diffusers
 - Diffusion Models Demystified

Part 5: Advanced Topics & Deployment (4 weeks)

- Topics:
 - RLHF (Reinforcement Learning with Human Feedback)
 - LangChain, LLM pipelines, vector stores (FAISS, Chroma)
 - LoRA, PEFT, quantization
 - Web deployment: Streamlit, Gradio, FastAPI
- Projects:
 - Chatbot with LangChain
 - Image captioning app
 - Web app for text/image generation

Capstone Ideas

- AI content creation suite (text, image, audio)
- Memory-enabled assistant (LangChain + ChromaDB)
- Custom AI-powered search system

Generative AI Learning Roadmap for Data Science Engineers

Bonus Resources

- The Illustrated Transformer: <https://jalammar.github.io/illustrated-transformer/>
- Full Stack Deep Learning: <https://fullstackdeeplearning.com/>
- Papers with Code Generative Models: <https://paperswithcode.com/task/generative-modeling>